

Lecture 1.2: Deep Feedforward Networks

Haizhao Yang

Department of Mathematics
University of Maryland College Park

2022 Summer Mini Course
Tianyuan Mathematical Center in Central China

- Overview of Deep Feedforward Networks
- Architecture Design
- Back-Propagation

Goal of Deep Feedforward Networks

Typical tasks of machine learning

- Learn a map from an input x to an output y ;

Goal of Deep Feedforward Networks

Typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;

Goal of Deep Feedforward Networks

Typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a finite family of maps $\{f(x; \theta)\}_{\theta}$;

Goal of Deep Feedforward Networks

Typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a finite family of maps $\{f(x; \theta)\}_\theta$;
- Create some criteria to quantify how good a map is;

Goal of Deep Feedforward Networks

Typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a finite family of maps $\{f(x; \theta)\}_\theta$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$

What's Deep Feedforward Networks?

Ideas of deep learning

- Instead of constructing the function approximation $f(x; \theta)$ directly, we use function composition to construct the approximation:

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ \dots \circ f_1^{\theta_1}(x) = f_L(f_{L-1}(\dots f_1(x; \theta_1) \dots; \theta_{L-1}); \theta_L),$$

where $\theta = (\theta_L, \dots, \theta_1)$.

What's Deep Feedforward Networks?

Ideas of deep learning

- Instead of constructing the function approximation $f(x; \theta)$ directly, we use function composition to construct the approximation:

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ \dots \circ f_1^{\theta_1}(x) = f_L(f_{L-1}(\dots f_1(x; \theta_1) \dots; \theta_{L-1}); \theta_L),$$

where $\theta = (\theta_L, \dots, \theta_1)$.

- $f_1^{\theta_1}$ is called the first layer;

What's Deep Feedforward Networks?

Ideas of deep learning

- Instead of constructing the function approximation $f(x; \theta)$ directly, we use function composition to construct the approximation:

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ \dots \circ f_1^{\theta_1}(x) = f_L(f_{L-1}(\dots f_1(x; \theta_1) \dots; \theta_{L-1}); \theta_L),$$

where $\theta = (\theta_L, \dots, \theta_1)$.

- $f_1^{\theta_1}$ is called the first layer;
- $f_L^{\theta_L}$ is called the last layer;

What's Deep Feedforward Networks?

Ideas of deep learning

- Instead of constructing the function approximation $f(x; \theta)$ directly, we use function composition to construct the approximation:

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ \dots \circ f_1^{\theta_1}(x) = f_L(f_{L-1}(\dots f_1(x; \theta_1) \dots; \theta_{L-1}); \theta_L),$$

where $\theta = (\theta_L, \dots, \theta_1)$.

- $f_1^{\theta_1}$ is called the first layer;
- $f_L^{\theta_L}$ is called the last layer;
- Other functions are called hidden layers and L is the depth of the model.

Questions for Deep Feedforward Networks

Question

- How to construct each layer $f(x, \theta_i)$?

Questions for Deep Feedforward Networks

Question

- How to construct each layer $f(x, \theta_i)$?
- What's the capacity of

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ f_1^{\theta_1}(x)?$$

Questions for Deep Feedforward Networks

Question

- How to construct each layer $f(x, \theta_i)$?
- What's the capacity of

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ f_1^{\theta_1}(x)?$$

- What's the efficiency of

$$f(x; \theta) = f_L^{\theta_L} \circ f_{L-1}^{\theta_{L-1}} \circ f_1^{\theta_1}(x)?$$

Models for each layer

- Linear models are nice and simple: $y = f(x; w) = w^T x$




Figure: SVM for classifying two set of points.

Models for each layer

- Linear models are nice and simple: $y = f(x; w) = w^T x$




Figure: SVM for classifying two set of points.

- However, nonlinear problems are more common




Figure: Need to introduce a nonlinear kernel mapping input x to its feature $\phi(x)$ such that we have a linear problem in the feature space. This is the kernel SVM $y = f(x; w, \phi) = w^T \phi(x)$.

Models for each layer




Figure: Kernel SVM for classifying two set of points.

Models for each layer

- Need a linear transform to shift points;
- Need a nonlinear transform to create linear separation;
- Need a linear transform for classification.




Figure: Kernel SVM for classifying two set of points.

Models for each layer

- Need a linear transform to shift points;
- Need a nonlinear transform to create linear separation;
- Need a linear transform for classification.




Figure: Kernel SVM for classifying two set of points.




Figure: In sum, we have motivated the simple 3 layer neuron network: $y = V^T \sigma(W^T x + b)$.

Models for each layer

- Below: FNN with 4 inputs, one hidden layer with 3 nodes, and 1 output




Figure: $y = V^T \sigma(W^T x + b)$.

Models for each layer

- Below: FNN with 4 inputs, one hidden layer with 3 nodes, and 1 output




Figure: $y = V^T \sigma(W^T x + b)$.

- Each hidden node computes a nonlinear transformation of its incoming inputs
 - Weighted linear combination followed by a nonlinear **activation function** σ

Models for each layer

- Below: FNN with 4 inputs, one hidden layer with 3 nodes, and 1 output




Figure: $y = V^T \sigma(W^T x + b)$.

- Each hidden node computes a nonlinear transformation of its incoming inputs
 - ▶ Weighted linear combination followed by a nonlinear **activation function** σ
 - ▶ Nonlinearity required by the nonlinear problem.

Models for each layer

- Below: FNN with 4 inputs, one hidden layer with 3 nodes, and 1 output





Figure: $y = V^T \sigma(W^T x + b)$.

- Each hidden node computes a nonlinear transformation of its incoming inputs
 - ▶ Weighted linear combination followed by a nonlinear **activation function** σ
 - ▶ Nonlinearity required by the nonlinear problem.
 - ▶ Output y is a linear transformation of the hidden nodes.


Models for each layer

- Below: A general 2 layer FNN



Models for each layer


- Below: A general 2 layer FNN



- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$

Models for each layer


- Below: A general 2 layer FNN



- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$

Models for each layer

- Below: A general 2 layer FNN




- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$, and a scalar-valued output node

$$y = \mathbf{v}^T \mathbf{h}$$

Models for each layer

- Below: A general 2 layer FNN




- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$, and a scalar-valued output node

$$y = \mathbf{v}^T \mathbf{h} = \mathbf{v}^T \sigma(\mathbf{W}^T \mathbf{x} + b),$$

where $\mathbf{v} = [v_1, \dots, v_K] \in \mathbb{R}^K$,

Models for each layer

- Below: A general 2 layer FNN




- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$, and a scalar-valued output node

$$y = \mathbf{v}^T \mathbf{h} = \mathbf{v}^T \sigma(\mathbf{W}^T \mathbf{x} + b),$$

where $\mathbf{v} = [v_1, \dots, v_K] \in \mathbb{R}^K$, $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K] \in \mathbb{R}^{D \times K}$,

Models for each layer

- Below: A general 2 layer FNN




- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$, and a scalar-valued output node

$$y = \mathbf{v}^T \mathbf{h} = \mathbf{v}^T \sigma(\mathbf{W}^T \mathbf{x} + b),$$

where $\mathbf{v} = [v_1, \dots, v_K] \in \mathbb{R}^K$, $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K] \in \mathbb{R}^{D \times K}$, σ is the nonlinear activation function.

Models for each layer

- Below: A general 2 layer FNN



- For an FNN with D inputs $\mathbf{x} = [x_1, \dots, x_D] \in \mathbb{R}^D$, a single layer with K hidden nodes $\mathbf{h} = [h_1, \dots, h_K]$, and a scalar-valued output node


$$y = \mathbf{v}^T \mathbf{h} = \mathbf{v}^T \sigma(\mathbf{W}^T \mathbf{x} + \mathbf{b}),$$

where $\mathbf{v} = [v_1, \dots, v_K] \in \mathbb{R}^K$, $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K] \in \mathbb{R}^{D \times K}$, σ is the nonlinear activation function.

- Each hidden node has a value $h_k = \sigma(\mathbf{w}_k^T \mathbf{x} + b_k) = \sigma(\sum_{d=1}^D w_{dk} x_d + b_k)$.

Deep Feedforward Neural Network

- Feed forward neural net with L hidden layers $\mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \dots, \mathbf{h}^{(L)}$, where $\mathbf{h}^{(1)} = \sigma(\mathbf{W}^{(1)T} \mathbf{x} + \mathbf{b}^{(1)})$ and $\mathbf{h}^{(\ell)} = \sigma(\mathbf{W}^{(\ell)T} \mathbf{h}^{(\ell-1)} + \mathbf{b}^{(\ell)})$ for $\ell \geq 2$.



- Note: The ℓ -th hidden layer contains K_ℓ hidden nodes, $\mathbf{W}^{(1)}$ is of size $D \times K_1$, $\mathbf{W}^{(\ell)}$ is of size $K_\ell \times K_{\ell+1}$, \mathbf{v} is of size $K_L \times 1$.

Why deep?

- Recall the kernel SVM




Figure: $y = f(x; w_1, w_2, b_1, b_2, \phi) = w_2^T \phi(w_1^T x + b_1) + b_2$.

Why deep?

- Recall the kernel SVM




Figure: $y = f(x; w_1, w_2, b_1, b_2, \phi) = w_2^T \phi(w_1^T x + b_1) + b_2$.

- Question: how to find the kernel $\phi(x)$ automatically?

Why deep?

- Recall the kernel SVM





Figure: $y = f(x; w_1, w_2, b_1, b_2, \phi) = w_2^T \phi(w_1^T x + b_1) + b_2$.

- Question: how to find the kernel $\phi(x)$ automatically?
- Use the composition of simple NN's with a simple nonlinearity $|x|$



Why deep?

- Recall the kernel SVM





Figure: $y = f(x; w_1, w_2, b_1, b_2, \phi) = w_2^T \phi(w_1^T x + b_1) + b_2$.

- Question: how to find the kernel $\phi(x)$ automatically?
- Use the composition of simple NN's with a simple nonlinearity $|x|$



- Identifying the best $\phi(x)$ is equivalent to optimizing the coefficients of deep NN (DNN).

Why deep? To be discussed later.

Open Problems for Research

- Approximation theory;
- Optimization;
- Generalization error.

- Overview of Deep Feedforward Networks
- Architecture Design
- Back-Propagation

Gradient-Based Learning

Recall typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a family of DNN's $\{f(x; \theta)\}_\theta$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$.

Gradient-Based Learning

Recall typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a family of DNN's $\{f(x; \theta)\}_\theta$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$.

Gradient-Based Learning

Recall typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a family of DNN's $\{f(x; \theta)\}_\theta$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$.

Principles

- The computation of gradient is fast;

Gradient-Based Learning

Recall typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a family of DNN's $\{f(x; \theta)\}_{\theta}$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$.

Principles

- The computation of gradient is fast;
- Gradient descent can quickly decrease the energy;

Gradient-Based Learning

Recall typical tasks of machine learning

- Learn a map from an input x to an output y ;
- Denote the map as $y = f(x)$;
- Construct a family of DNN's $\{f(x; \theta)\}_{\theta}$;
- Create some criteria to quantify how good a map is;
- Use optimization to find the best $f(x; \theta)$.

Principles

- The computation of gradient is fast;
- Gradient descent can quickly decrease the energy;
- Local minimizer is reasonably good;

Typical cost functions

Learning conditional distributions

- We often maximize the likelihood function to obtain a distribution that best matches given data:

$$J(\theta) = -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}^{\theta}(\mathbf{y}|\mathbf{x}).$$

Typical cost functions

Learning conditional distributions

- We often maximize the likelihood function to obtain a distribution that best matches given data:

$$J(\theta) = -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}^{\theta}(\mathbf{y}|\mathbf{x}).$$

- Specifying a model $p^{\theta}(\mathbf{y}|\mathbf{x})$ automatically determines a cost function.

Typical cost functions

Learning conditional distributions

- We often maximize the likelihood function to obtain a distribution that best matches given data:

$$J(\theta) = -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}^{\theta}(\mathbf{y}|\mathbf{x}).$$

- Specifying a model $p^{\theta}(\mathbf{y}|\mathbf{x})$ automatically determines a cost function.
- When we use an DNN to specify the distribution function, say

$$p^{\theta}(\mathbf{y}|\mathbf{x}) = f(\mathbf{y}; \mathbf{x}, \theta),$$

we need to choose the output and hidden units carefully so that the objective function is easy to optimize.

Typical cost functions

Example 1: Mean squared error

- Assumption: $y = f(\mathbf{x}; \theta) + \omega$, where ω is a Gaussian random noise.
- Model: $p_{model}^{\theta}(\mathbf{y}|\mathbf{x}) = \exp(-\frac{\|\mathbf{y} - f(\mathbf{x}; \theta)\|_2^2}{2\sigma^2})$.
- Loss function in the MLE framework:

$$\begin{aligned}
 J(\theta) &= -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}^{\theta}(\mathbf{y}|\mathbf{x}) \\
 &= \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \frac{\|\mathbf{y} - f(\mathbf{x}; \theta)\|_2^2}{2\sigma^2} \\
 &= \frac{1}{N} \sum_{i=1}^N \frac{\|\mathbf{y}_i - f(\mathbf{x}_i; \theta)\|_2^2}{2\sigma^2}
 \end{aligned}$$

- Optimization problem:

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \|\mathbf{y}_i - f(\mathbf{x}_i; \theta)\|_2^2$$

Typical cost functions

Example 2: Mean absolute error

- Assumption: $y = f(\mathbf{x}; \theta) + \omega$, where ω is a Laplace random noise.
- Model: $p_{model}^{\theta}(\mathbf{y}|\mathbf{x}) = \frac{1}{2b} \exp(-\frac{\|\mathbf{y} - f(\mathbf{x}; \theta)\|_1}{b})$.
- Loss function in the MLE framework:

$$\begin{aligned}
 J(\theta) &= -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}^{\theta}(\mathbf{y}|\mathbf{x}) \\
 &= \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \frac{\|\mathbf{y} - f(\mathbf{x}; \theta)\|_1}{b} + \log(2b) \\
 &= \frac{1}{N} \sum_{i=1}^N \frac{\|\mathbf{y}_i - f(\mathbf{x}_i; \theta)\|_1}{b} + \log(2b)
 \end{aligned}$$

- Optimization problem:

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \|\mathbf{y}_i - f(\mathbf{x}_i; \theta)\|_1$$

Typical cost functions

Example 3: Mean squared log error

- Assumption: $\log y = \log f(\mathbf{x}; \theta) + \omega$, where ω is a Gaussian random noise.
- Model: $p_{model}^{\theta}(\log \mathbf{y} | \mathbf{x}) = \exp(-\frac{\|\log \mathbf{y} - \log f(\mathbf{x}; \theta)\|_2^2}{2\sigma^2})$.
- Loss function in the MLE framework:

$$\begin{aligned}
 J(\theta) &= -\mathbb{E}_{\mathbf{x}, \log y \sim \hat{p}_{data}} \log p_{model}^{\theta}(\log \mathbf{y} | \mathbf{x}) \\
 &= \mathbb{E}_{\mathbf{x}, \log y \sim \hat{p}_{data}} \frac{\|\log \mathbf{y} - \log f(\mathbf{x}; \theta)\|_2^2}{2\sigma^2} \\
 &= \frac{1}{N} \sum_{i=1}^N \frac{\|\log \mathbf{y}_i - \log f(\mathbf{x}_i; \theta)\|_2^2}{2\sigma^2}
 \end{aligned}$$

- Optimization problem:

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \|\log \mathbf{y}_i - \log f(\mathbf{x}_i; \theta)\|_2^2$$

Typical cost functions


MSE vs. MSLE

We use MSLE instead of MSE when

- when you don't want to penalize huge differences in the predicted and the actual values when both predicted and true values are huge numbers.
- when you want to penalize under estimates more than over estimates. For example: P_i as predicted value, A_i as actual value,
 - ▶ $P_i = 600$, $A_i = 1000$, then
 $RMSE = 400$, $RMSLE = 0.5108$
 - ▶ $P_i = 1400$, $A_i = 1000$, then
 $RMSE = 400$, $RMSLE = 0.3365$

Typical cost functions

Other loss functions



- Cross entropy (classification);
- KL divergence (distribution);
- EMD (distribution);
- Huber (regression);
- (Squared) Hinge (classification).


Reference: On Loss Functions for Deep Neural Networks in Classification
<https://arxiv.org/abs/1702.05659>.

Choices of Units

- The choice of cost function is tightly coupled with the choice of output and hidden units


Choices of Units

- The choice of cost function is tightly coupled with the choice of output and hidden units
- Nonlinear activation functions:
 - Sigmoid: $f(x) = \sigma(x) = \frac{1}{1+\exp(-x)}$ (range between 0-1)
 - tanh: $f(x) = 2\sigma(2x) - 1$ (range between -1 and +1)
 - Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$



Choices of Units

- The choice of cost function is tightly coupled with the choice of output and hidden units
- Nonlinear activation functions:
 - Sigmoid: $f(x) = \sigma(x) = \frac{1}{1+\exp(-x)}$ (range between 0-1)
 - tanh: $f(x) = 2\sigma(2x) - 1$ (range between -1 and +1)
 - Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$




- Sigmoid saturates and can kill gradients

Choices of Units

- The choice of cost function is tightly coupled with the choice of output and hidden units
- Nonlinear activation functions:

- Sigmoid: $f(x) = \sigma(x) = \frac{1}{1 + \exp(-x)}$ (range between 0-1)
- tanh: $f(x) = 2\sigma(2x) - 1$ (range between -1 and +1)
- Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$




- Sigmoid saturates and can kill gradients
- tanh also saturates but is steeper around the center (thus preferred over sigmoid)

Choices of Units

- The choice of cost function is tightly coupled with the choice of output and hidden units
- Nonlinear activation functions:

- Sigmoid: $f(x) = \sigma(x) = \frac{1}{1+\exp(-x)}$ (range between 0-1)
- tanh: $f(x) = 2\sigma(2x) - 1$ (range between -1 and +1)
- Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$




- Sigmoid saturates and can kill gradients
- tanh also saturates but is steeper around the center (thus preferred over sigmoid)
- ReLU is currently the most popular (also cheap to compute), leading to piecewise linear functions.

- Overview of Deep Feedforward Networks
- Architecture Design
- Back-Propagation

Back-Propagation


- Want to learn the parameters by minimizing some loss function



- Backpropagation** (gradient descent + chain rule for derivatives) is commonly used to do this efficiently

Back-Propagation

- Consider the feedforward neural net with one hidden layer



- Recall that $\mathbf{h} = [h_1 \ h_2 \ \dots \ h_K] = f(\mathbf{W}^T \mathbf{x})$
- Assuming a regression problem, the optimization problem would be

$$\min_{\mathbf{W}, \mathbf{v}} \frac{1}{2} \sum_{n=1}^N \left(y_n - \mathbf{v}^T f(\mathbf{W}^T \mathbf{x}_n) \right)^2 = \min_{\mathbf{W}, \mathbf{v}} \frac{1}{2} \sum_{n=1}^N \left(y_n - \sum_{k=1}^K v_k f(\mathbf{w}_k^T \mathbf{x}_n) \right)^2$$

where \mathbf{w}_k is the k -th column of the $D \times K$ matrix \mathbf{W}

Back-Propagation

- We can learn the parameters by doing gradient descent (or stochastic gradient descent) on the objective function

$$\mathcal{L} = \frac{1}{2} \sum_{n=1}^N \left(y_n - \sum_{k=1}^K v_k f(\mathbf{w}_k^\top \mathbf{x}_n) \right)^2 = \frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{v}^\top \mathbf{h}_n)^2$$

- Gradient w.r.t. $\mathbf{v} = [v_1 \ v_2 \ \dots \ v_K]$ is straightforward

$$\frac{\partial \mathcal{L}}{\partial \mathbf{v}} = - \sum_{n=1}^N \left(y_n - \sum_{k=1}^K v_k f(\mathbf{w}_k^\top \mathbf{x}_n) \right) \mathbf{h}_n = - \sum_{n=1}^N \mathbf{e}_n \mathbf{h}_n$$


- Gradient w.r.t. the weights $\mathbf{W} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_K]$ is a bit more involved due to the presence of f but can be computed using chain rule

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}_k} = \frac{\partial \mathcal{L}}{\partial f_k} \frac{\partial f_k}{\partial \mathbf{w}_k} \quad (\text{note: } f_k = f(\mathbf{w}_k^\top \mathbf{x}))$$

- We have: $\frac{\partial \mathcal{L}}{\partial f_k} = - \sum_{n=1}^N (y_n - \sum_{k=1}^K v_k f(\mathbf{w}_k^\top \mathbf{x}_n)) v_k = - \sum_{n=1}^N \mathbf{e}_n v_k$
- We have: $\frac{\partial f_k}{\partial \mathbf{w}_k} = \sum_{n=1}^N f'(\mathbf{w}_k^\top \mathbf{x}_n) \mathbf{x}_n$, where $f'(\mathbf{w}_k^\top \mathbf{x}_n)$ is f 's derivative at $\mathbf{w}_k^\top \mathbf{x}_n$
- These calculations can be done efficiently using [backpropagation](#)

Back-Propagation

- Basically consists of a forward pass and a backward pass



- Forward pass computes the errors e_n using the current parameters
- Backward pass computes the gradients and updates the parameters, starting from the parameters at the top layer and then moving backwards
- Also good at reusing previous computations (updates of parameters at any layer depends on parameters at the layer above)